

# HelpMate AI

Project Report by Sandeep B.



## Objectives



### **Develop a Semantic Search System**

Utilize the RAG pipeline for efficient document retrieval, integrating embedding, search, and generation layers.



### **Implement a Cache Layer**

Enhance performance by caching previous queries and results using ChromaDB.



### **Extract and Structure Information**

Process and store PDF documents in a structured format, generating vector representations with OpenAI's models.



### **Build a Generative Search System**

Develop a robust system capable of accurately answering questions from a policy document.



# System Design

## RAG Pipeline & Cache Implementation



### **RAG Pipeline Overview**

Integrates Embedding, Search, and Generation layers to retrieve and generate accurate responses.



### **Search & Rank Layer**

Performs semantic similarity searches and re-ranks results using cross-encoders.



### **Embedding Layer**

Processes and chunks PDFs, generating vector embeddings stored in ChromaDB.



### **Cache Implementation**

Enhances performance by caching queries and results, utilizing a similarity threshold.



## Implementation

- **Document Processing:** Utilized pdfplumber for text and table extraction, followed by chunking and vector embedding with OpenAI's models.
- **Semantic Search:** Implemented semantic similarity searches using the RAG pipeline and vector database ChromaDB.
- **Cache System:** Developed a cache system with ChromaDB, optimizing the retrieval of previous queries and results.



Photo by Maik Jonietz on Unsplash

## Challenges Faced

- **Data Quality & Preprocessing:** Extracting relevant information from complex insurance documents proved challenging due to varied text structures.
- **Chunking Strategies:** Optimizing chunk size and overlap to maintain context without losing coherence was critical but difficult.
- **Query Understanding & Matching:** Designing relevant queries that required sophisticated understanding and reasoning posed a significant challenge.

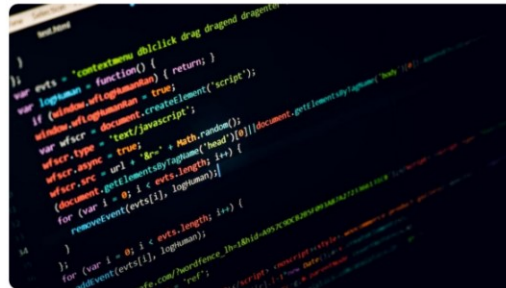


Photo by RoonZ nl on Unsplash

## Conclusion



### **Successful Implementation**

HelpMate AI successfully achieved its objectives, implementing a robust semantic search system with the RAG pipeline.



### **Challenges Overcome**

The project addressed significant challenges in data processing, chunking strategies, and query design.



### **Scalable and Efficient System**

The final system is scalable, efficient, and provides accurate information retrieval from complex documents.

---

## Lessons Learned



### **Efficient Document Processing**

Utilizing tools like pdfplumber is crucial for handling complex PDF documents efficiently.



### **Semantic Search Optimization**

Fine-tuning search parameters and thresholds is essential for achieving optimal results.



### **Cache Management**

Implementing an effective caching strategy significantly improves system performance.

